

# Analyzing volatility

GARCH MODELS IN R



**Kris Boudt**

Professor of finance and econometrics

# About the instructor

- Kris Boudt
  - PhD in financial risk forecasting
  - Use GARCH models to win by not losing (much)
- R package rugarch of Alexios Ghalanos.



# Calculating returns

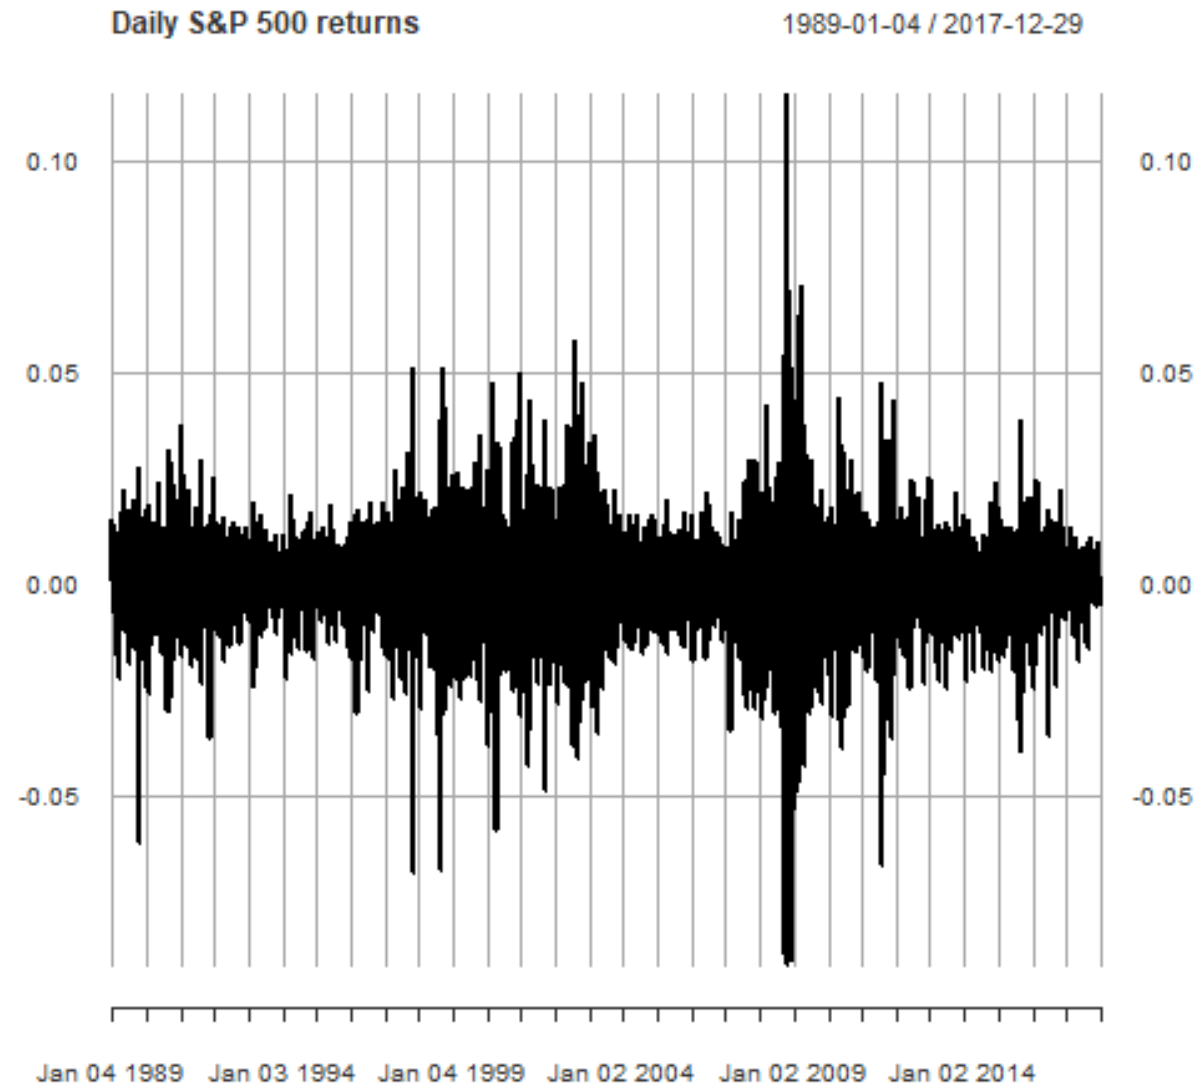
- Relative financial gains and losses, expressed in terms of returns

$$R_t = \frac{P_t - P_{t-1}}{P_{t-1}}$$

- Function `CalculateReturns` in `PerformanceAnalytics`

```
# Example in R for daily S&P 500 prices (xts object)
library(PerformanceAnalytics)
SP500returns <- CalculateReturns(SP500prices)
```

# Daily S&P 500 returns




Properties of daily returns:

- The average return is zero
- Return variability changes through time

Standard deviation = measure of return variability.

Synonym: Return **volatility**.

Greek letter  $\sigma_t$ .



# How to estimate return volatility

- Function `sd()` computes the (daily) standard deviation:

```
sd(sp500ret)
```

```
0.01099357
```

- Corresponding formula for  $\hat{\sigma}$  computed using  $T$  daily returns:

$$\hat{\sigma} = \sqrt{\frac{1}{T-1} \sum_{t=1}^T (R_t - \hat{\mu})^2},$$

- where  $\hat{\mu}$  is the mean return.

# Annualized volatility


- `sd(sp500ret)` is daily volatility
- Annualized volatility =  $\sqrt{252} \times$  daily volatility

```
# Compute annualized standard deviation  
sqrt(252) * sd(sp500ret)
```

```
0.1745175
```

S&P 500 returns in 1989

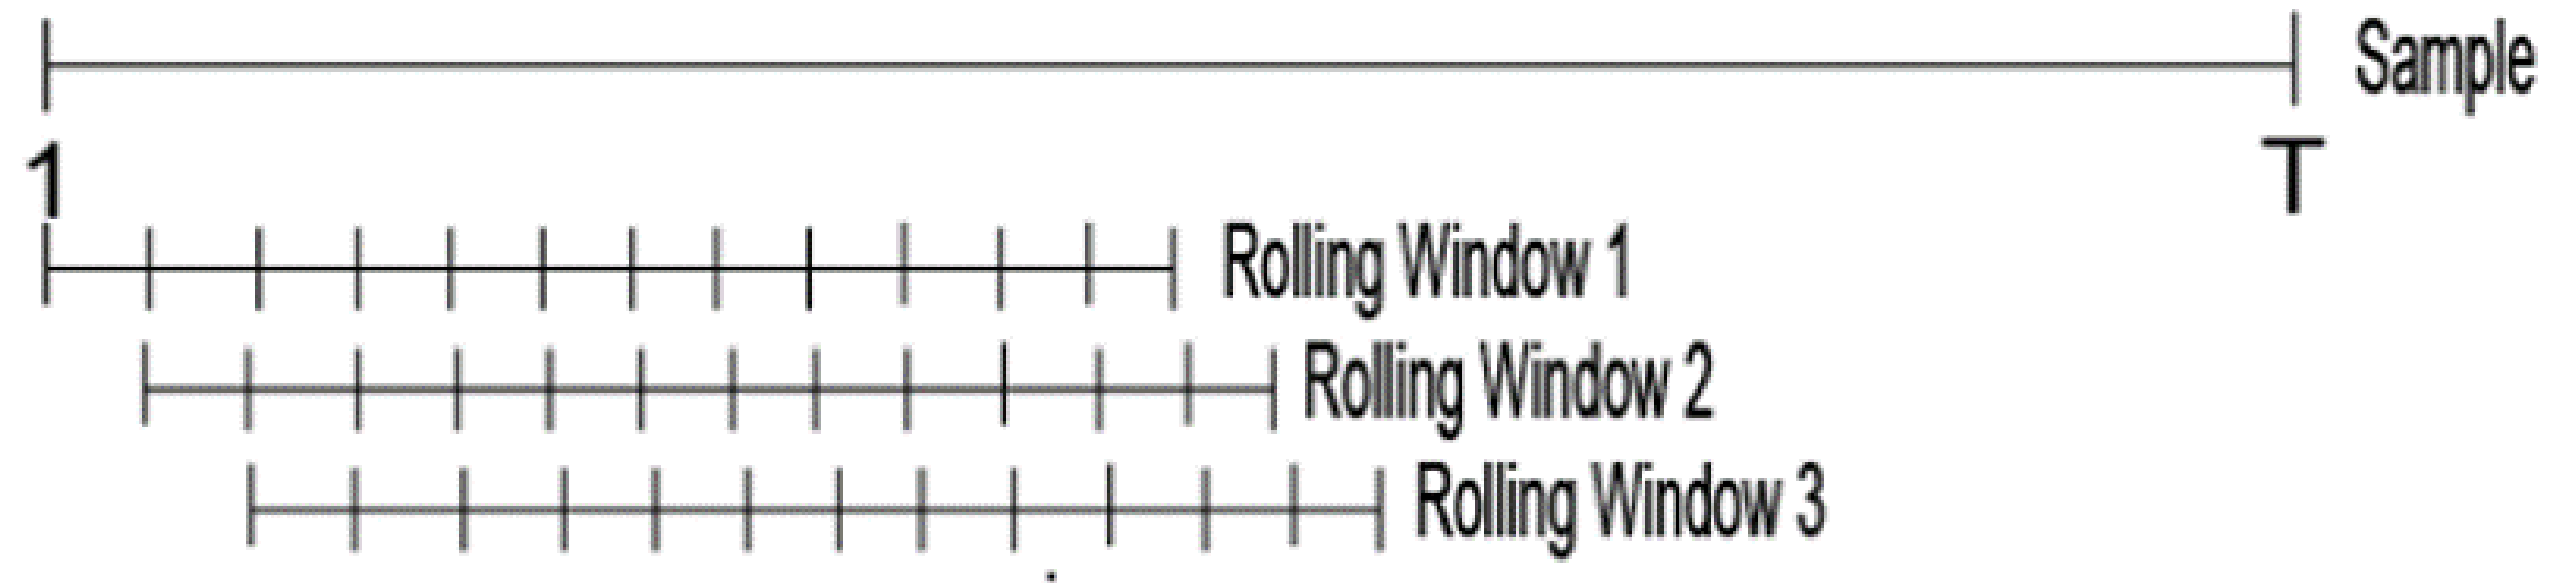
Ann. volatility=13%





# Rolling volatility estimation

- Rolling estimation windows :



- Window width? Multiple of 22 (trading days).


# Function `chart.RollingPerformance()`

```
library(PerformanceAnalytics)

chart.RollingPerformance(R = sp500ret,
                        width = 22,
                        FUN = "sd.annualized",
                        scale = 252,
                        main = "Rolling 1 month volatility")
```

Rolling volatility

1989-01-04 / 2017-12-29



# About GARCH models in R

- Estimation of  $\sigma_t$  requires time series models, like GARCH.

# Let's refresh the basics of computing rolling standard deviations in R

GARCH MODELS IN R

# GARCH models: The way forward

GARCH MODELS IN R



**Kris Boudt**

Professor of finance and econometrics

# Inventors of GARCH models

Robert Engle




Tim Bollerslev



# Notation (i)


- Input: Time series of returns






# Notation (ii)

- At time  $t - 1$ , you make the prediction about the the future return  $R_t$ , using the *information set* available at time  $t - 1$ :



# Notation (iii)

- Predicting the **mean return**: what is the best possible prediction of the actual return?




$I_{t-1}$  = Information set available at the time of prediction ( $t-1$ )

$$\mu_t = E[R_t \mid I_{t-1}]$$

Prediction error:  $e_t = R_t - \mu_t$

# Notation (iv)

- We then predict the **variance**: how far off the return can be from its mean?




$I_{t-1}$  = Information set available at the time of prediction (t-1)

$$\begin{aligned}\sigma_t^2 &= \text{var}(R_t \mid I_{t-1}) \\ &= E[(R_t - \mu_t)^2 \mid I_{t-1}] \\ &= E[e_t^2 \mid I_{t-1}]\end{aligned}$$

$$\sigma_t = \sqrt{\sigma_t^2}$$

# From theory to practice: Models for the mean

- We need an equation that maps the past returns into a prediction of the mean




$$\text{Rolling mean model: } \mu_t = \frac{1}{M} \sum_{i=1}^M R_{t-i}$$

For AR(MA) models for the mean, see Datacamp course on [time series analysis](#).

# From theory to practice: Models for the variance


- We need an equation that maps the past returns into predictions of the variance



Rolling variance model:  $\sigma_t^2 = \frac{1}{M} \sum_{i=1}^M e_{t-i}^2$

# ARCH(p) model: Autoregressive Conditional Heteroscedasticity

- We need an equation that maps the past returns into predictions of the variance




Rolling variance model:  $\sigma_t^2 = \frac{1}{M} \sum_{i=1}^M e_{t-i}^2$

ARCH(p) model:  $\sigma_t^2 = \omega + \sum_{i=1}^p \alpha_i e_{t-i}^2$

# GARCH(1,1) model: Generalized ARCH

- We need an equation that maps the past returns into predictions of the variance



$I_{t-1}$  = Information set available at the time of prediction (t-1)

$$\text{ARCH}(p) \text{ model: } \sigma_t^2 = \omega + \sum_{i=1}^p \alpha_i e_{t-i}^2$$

$$\text{GARCH}(1,1) \text{ model: } \sigma_t^2 = \omega + \alpha e_{t-1}^2 + \beta \sigma_{t-1}^2$$

# Parameter restrictions

To make the GARCH process realistic, we need that:

1.  $\omega$ ,  $\alpha$  and  $\beta$  are  $> 0$ : this ensures that  $\sigma_t^2 > 0$  at all times.
2.  $\alpha + \beta < 1$ : this ensures that the predicted variance  $\sigma_t^2$  always returns to the long run variance:
  - The variance is therefore "mean-reverting"
  - The long run variance equals  $\frac{\omega}{1-\alpha-\beta}$



# R implementation - Specify the inputs

- Let's familiarize ourselves with the GARCH equations using R code:

$$\sigma_t^2 = \omega + \alpha e_{t-1}^2 \beta \sigma_{t-1}^2$$

```
# Set parameter values
alpha <- 0.1
beta <- 0.8
omega <- var(sp500ret) * (1 - alpha - beta)
# Then: var(sp500ret) = omega / (1 - alpha - beta)
```

```
# Set series of prediction error
e <- sp500ret - mean(sp500ret) # Constant mean
e2 <- e ^ 2
```

# R implementation - compute predicted variances

```
# We predict for each observation its variance.
```

```
nobs <- length(sp500ret)
```

```
predvar <- rep(NA, nobs)
```

```
# Initialize the process at the sample variance
```

```
predvar[1] <- var(sp500ret)
```

```
# Loop starting at 2 because of the lagged predictor
```

```
for (t in 2:nobs){
```

```
  predvar[t] <- omega + alpha * e2[t - 1] + beta * predvar[t-1]
```

```
}
```

# R implementation - Plot of GARCH volatilities


```
# Volatility is sqrt of predicted variance
predvol <- sqrt(predvar)
predvol <- xts(predvol, order.by = time(sp500ret))
```

```
# We compare with the unconditional volatility
uncvol <- sqrt(omega / (1 - alpha-beta))
uncvol <- xts(rep(uncvol, nobs), order.by = time(sp500ret))
```

```
# Plot
plot(predvol)
lines(uncvol, col = "red", lwd = 2)
```

predvol

1989-01-04 / 2017-12-29



# Let's practice!

GARCH MODELS IN R

# The rugarch package

GARCH MODELS IN R



**Kris Boudt**

Professor of finance and econometrics

# The normal GARCH(1,1) model with constant mean

The normal GARCH model

$$R_t = \mu + e_t$$

$$e_t \sim N(0, \sigma_t^2)$$

$$\sigma_t^2 = \omega + \alpha e_{t-1}^2 + \beta \sigma_{t-1}^2$$

- Four parameters:  $\mu, \omega, \alpha, \beta$ .
- Estimation by **maximum likelihood**: find the parameter values for which the GARCH model is most likely to have generated the observed return series.

# Alexios Ghalanos

```
library(rugarch)  
citation("rugarch")
```

To cite the rugarch package, please use:

Alexios Ghalanos (2018). rugarch: Univariate GARCH models.R package version 1.4-0.



# Workflow

- Three steps:
  - `ugarchspec()` : Specify which GARCH model you want to use (mean  $\mu_t$ , variance  $\sigma_t^2$ , distribution of  $e_t$ )
  - `ugarchfit()` : Estimate the GARCH model on your time series with returns  $R_1, \dots, R_T$ .
  - `ugarchforecast()` : Use the estimated GARCH model to make volatility predictions for  $R_{T+1}, \dots$

# Workflow in R

`ugarchspec()` specifies which GARCH model you want to use.

```
garchspec <- ugarchspec(mean.model = list(armaOrder = c(0,0)),  
                        variance.model = list(model = "sGARCH"),  
                        distribution.model = "norm")
```

`ugarchfit()` estimates the GARCH model.

```
garchfit <- ugarchfit(data = sp500ret, spec = garchspec)
```

`ugarchforecast()` forecasts the volatility of the future returns.

```
garchforecast <- ugarchforecast(fit0Rspec = garchfit, n.ahead = 5)
```

# ugarchfit object

- The `ugarchfit` yields an object that contains all the results related to the estimation of the garch model.
- Methods `coef`, `uncvar`, `fitted` and `sigma`:

```
# Coefficients
garchcoef <- coef(garchfit)
# Unconditional variance
garchuncvar <- uncvariance(garchfit)
# Predicted mean
garchmean <- fitted(garchfit)
# Predicted volatilities
garchvol <- sigma(garchfit)
```

# GARCH coefficients for daily S&P 500 returns

```
print(garchcoef)
```

```
      mu      omega    alpha1    beta1  
5.728020e-04 1.220515e-06 7.792031e-02 9.111455e-01
```

$$R_t = 5.7 \times 10^{-4} + e_t$$

$$e_t \sim N(0, \hat{\sigma}_t^2)$$


$$\hat{\sigma}_t^2 = 1.2 \times 10^{-6} + 0.08e_{t-1}^2 + 0.91\hat{\sigma}_{t-1}^2$$

```
sqrt(garchuncvar)
```

```
0.01056519
```

# Estimated volatilities

```
garchvol <- sigma(garchfit)  
plot(garchvol)
```



# What about future volatility?

```
tail(garchvol, 1)
```

```
2017-12-29 0.004862908
```

What about the volatility for the days following the end of the time series?

# Forecasting h-day ahead volatilities

- Applying the `sigma()` method to the `ugarchforecast` object gives the volatility forecasts:

```
sigma(garchforecast)
```

```
      2017-12-29
T+1  0.005034754
T+2  0.005127582
T+3  0.005217770
T+4  0.005305465
T+5  0.005390797
```

# Forecasting h-day ahead volatilities

Applying the `fitted()` method to the `ugarchforecast` object gives the mean forecasts:

```
fitted(garchforecast)
```

```
      2017-12-29
T+1  0.000572802
T+2  0.000572802
T+3  0.000572802
T+4  0.000572802
T+5  0.000572802
```



# Application to tactical asset allocation

A portfolio that invests a percentage  $w$  in a risky asset (with volatility  $\sigma_t$ ) and keeps  $1 - w$  on a risk-free bank deposit account has volatility equal to

$$\sigma_p = w\sigma_t$$

How to set  $w$ ? One approach is **volatility targeting**:  $w$  is such that the predicted annualized portfolio volatility equals a target level, say 5%. Then:

$$w^* = 0.05 / \sigma_t$$

Since GARCH volatilities change, the optimal weight changes as well.

# Let's play with rugarch!

GARCH MODELS IN R